

WENHU CHEN

Email: (hustchenwenhu@gmail.com, wenhuchen@uwaterloo.ca), **Website:** <https://wenhuchen.github.io/>

ABOUT ME

I'm currently an LM researcher at Meta SuperIntelligence Labs. I was a tenure-track assistant professor at Computer Science of University of Waterloo and Vector Institute since Aug 2022. I became a Canadian CIFAR AI chair since Jan 2023. I mainly work on different aspects of foundation models. My research can be summarized as the following aspects:

- 1) **large language models and agents:** post-training, reasoning, coding, agents, etc.
- 2) **multimodal language models:** post-training, multi-image/video understanding, etc.
- 3) **retrieval augmentation:** information retrieval, retrieval-augmented generation, etc.
- 4) **evaluation:** benchmarks, automatic metrics, reward models, etc.

WORKING EXPERIENCE

Meta SuperIntelligence Labs

02/13/2025-Now

Location: Menlo Park, US

Role: Research Scientist

Duties: 1. Research on Multimodal Data and Evaluation. 2. Research on Pre-training Science.

Perplexity AI

11/03/2025-02/09/2026

Location: San Francisco, US

Role: Member of Technical Staff

Duties: 1. Building search agents to perform high-accuracy and precision information seeking. 2. Post-training open-source models to match ChatGPT on human preferences (Perplexity Sonar).

Vector Institute

08/22/2022-Now

Location: Toronto, Canada

Role: Faculty Member (On Leave)

Duties: 1. Conducting research in natural language processing. 2. Service in postdoc and faculty review committee. 3. Hosting research events in Vector Institute.

University of Waterloo

08/22/2022-Now

Location: Waterloo, Canada

Role: Assistant Professor in Cheriton School of Computer Science (On Leave)

Duties: 1. Leading research projects and supervising graduate students. 2. Teaching fundamental courses. 3. Administrative services (i.e. the graduate admission committee, undergrad hiring committee).

Google Deepmind (Gemini Team)

08/22/2022-07/28/2025

Location: Remote in Canada

Role: Senior Research Scientist (20% Part-Time)

Manager: William W. Cohen & Ming-Wei Chang

Duties: Contributed to Gemini Multimodal and Evaluation efforts. I mainly participated in the Gemini 2.5 and the Nano-Banana project.

Google Research

08/16/2021-08/19/2022

Location: Pittsburgh, USA

Role: Research Scientist (Full time)

Manager: William W. Cohen

Duties: Building Next-Generation Retrieval-Augmented Systems to enhance factuality.

EDUCATION

University of California, Santa Barbara
Location: Santa Barbara, USA
Ph.D in Computer Science
Advisor: William (Yang) Wang and Xifeng Yan

09/01/2017-06/15/2021

RWTH Aachen University, Germany
Location: Aachen, Germany
M.Sc. in Electronics Engineering

09/01/2014-02/01/2017

Huazhong University of Science and Technology, China
Location: Wuhan, China
B.Sc. in Electronics Engineering

09/01/2010-06/15/2014

INTERNSHIP EXPERIENCE

Google Research June 15th, 2020 - Dec 18th, 2020
Location: Pittsburgh, PA, USA
Mentors: Ming-Wei Chang, Eva Schlinger, William W. Cohen
Duties: Work as research intern to explore research topics. We Build question answering system that can understand both structured and unstructured information. Our work was published in [ICLR 2021].

Microsoft AI & Cloud June 17th, 2019 - Sep 6th, 2019
Location: Seattle, WA, USA
Mentors: Zhe Gan, Jingjing Liu
Duties: Work as research intern to explore research topics. We build explainable visual question answering system that can ground on images to answer simple questions. Our work was published in [WACV 2021].

Samsung Research America June 18th, 2018 - Sep 14th, 2018
Location: Mountain View, CA, USA
Mentors: Yi-Lin Shen, Hongxia Jin
Duties: Work as research intern to explore research topics. We investigate different compression algorithm in natural language systems to save computation and storage. Our work was published in [NAACL 2019].

Microsoft Research Asia Feb 13th, 2017 - Aug 11th, 2017
Location: Beijing, China
Mentors: Shujie Liu, Ming Zhou
Duties: Work as research intern to explore research topics. We enhance the machine translation systems by utilizing data augmentation strategies. Our work was published in [NAACL 2018].

eBay Research Dec 7th, 2015 - May 27th, 2016
Location: Aachen, Germany
Mentors: Matusov Evgency, Shahram Khadivi
Duties: Work as research intern to explore research topics. We enhance machine translation with statistical alignment algorithm. Our work was published in [AMTA 2016].

TEACHING EXPERIENCE

University of Waterloo, 2023 Winter : CS 486/686 Introduction to Artificial Intelligence
University of Waterloo, 2024 Winter : CS 886 Recent Advances on Foundation Models
University of Waterloo, 2024 Spring : CS 486/686 Introduction to Artificial Intelligence

SELECTED PUBLICATIONS (AS FIRST & LAST AUTHOR)

[TMLR 2026] VideoEval-Pro: Robust and Realistic Long Video Understanding Evaluation
Wentao Ma, Weiming Ren, Yiming Jia, Zhuofeng Li, Ping Nie, Ge Zhang, Wenhui Chen

[ICML 2026] Context Forcing: Consistent Autoregressive Video Generation with Long Context
Shuo Chen, Cong Wei, Sun Sun, Ping Nie, Kai Zhou, Ge Zhang, Ming-Hsuan Yang, Wenhui Chen

[**TMLR 2026**] VerlTool: Towards Holistic Agentic Reinforcement Learning with Tool Use
Dongfu Jiang, Yi Lu, Zhuofeng Li, Zhiheng Lyu, Ping Nie, Haozhe Wang, Alex Su, Hui Chen, Kai Zou, Chao Du, Tianyu Pang, Wenhu Chen
(**Best Paper Award on SPOT@ICLR26**)

[**Findings of ACL 2026**] SWE-QA-Pro: A Representative Benchmark and Scalable Training Recipe for Repository-Level Code Understanding
Songcheng Cai, Zhiheng Lyu, Yuansheng Ni, Xiangchao Chen, Baichuan Zhou, Shenzhe Zhu, Yi Lu, Haozhe Wang, Chi Ruan, Benjamin Schneider, Weixu Zhang, Xiang Li, Andy Zheng, Yuyu Zhang, Ping Nie, Wenhu Chen

[**ICLR 2026**] ImagenWorld: Stress-Testing Image Generation Models with Explainable Human Evaluation on Open-ended Real-World Tasks
Samin.Mahdizadeh.Sani1, Max Ku, Nima Jamali, Matina Mahdizadeh Sani, Paria Khoshtab, Wei-Chieh Sun, Parnian Fazel, Zhi Rui Tam, Thomas Chong, Edisy Kin Wai Chan, Donald Wai Tong Tsang, Chiao-Wei Hsu, Lam Ting Wai, Ho Yin Sam Ng, Chiafeng Chu, Chak-Wing Mak, Keming Wu, Hiu Tung Wong, Yik Chun Ho, Chi Ruan, Zhuofeng Li, I-Sheng Fang, Shih-Ying Yeh, Ho Kei Cheng, Ping Nie, Wenhu Chen

[**ICLR 2026**] Critique-Coder: Enhancing Coder Models by Critique Reinforcement Learning
Chi Ruan, Dongfu Jiang, Yubo Wang, Wenhu Chen

[**ICLR 2026**] VisCoder2: Building Multi-Language Visualization Coding Agents
Yuansheng Ni, Songcheng Cai, Xiangchao Chen, Jiarong Liang, Zhiheng Lyu, Jiaqi Deng, Kai Zou, Ping Nie, Fei Yuan, Xiang Yue, Wenhu Chen

[**ICLR 2026**] UniVideo: Unified Understanding, Generation, and Editing for Videos
Cong Wei, Quande Liu, Zixuan Ye, Qiulin Wang, Xintao Wang, Pengfei Wan, Kun Gai, Wenhu Chen

[**ICLR 2026**] EditReward: A Human-Aligned Reward Model for Instruction-Guided Image Editing
Keming Wu, Sicong Jiang, Max Ku, Ping Nie, Minghao Liu, Wenhu Chen

[**ICLR 2026**] Emergent Hierarchical Reasoning in LLMs through Reinforcement Learning
Haozhe Wang, Qixin Xu, Che Liu, Junhong Wu, Fangzhen Lin, Wenhu Chen

[**TMLR 2025**] BrowserAgent: Building Web Agents with Human-Inspired Web Browsing Actions
Tao Yu, Zhengbo Zhang, Zhiheng Lyu, Junhao Gong, Hongzhu Yi, Xinming Wang, Yuxuan Zhou, Jiabing Yang, Ping Nie, Yan Huang, Wenhu Chen

[**TMLR 2025**] QuickVideo: Real-Time Long Video Understanding with System Algorithm Co-Design
Benjamin Schneider, Dongfu Jiang, Chao Du, Tianyu Pang, Wenhu Chen

[**TMLR 2025**] StructEval: Benchmarking LLMs' Capabilities to Generate Structural Outputs
Jialin Yang, Dongfu Jiang, Lipeng He, Sherman Siu, Yuxuan Zhang, Disen Liao, Zhuofeng Li, Huaye Zeng, Yiming Jia, Haozhe Wang, Benjamin Schneider, Chi Ruan, Wentao Ma, Zhiheng Lyu, Yifei Wang, Yi Lu, Quy Duc Do, Ziyang Jiang, Ping Nie, Wenhu Chen

[**NeurIPS 2025**] Pixel Reasoner: Incentivizing Pixel-Space Reasoning with Curiosity-Driven Reinforcement Learning
Alex Su, Haozhe Wang, Weimin Ren, Fangzhen Lin, Wenhu Chen

[**NeurIPS 2025**] General-Reasoner: Advancing LLM Reasoning Across All Domains
Xueguang Ma, Qian Liu, Dongfu Jiang, Ge Zhang, Zejun Ma, Wenhu Chen

[**NeurIPS 2025**] VL-Rethinker: Incentivizing Self-Reflection of Vision-Language Models with Reinforcement Learning
Haozhe Wang, Chao Qu, Zuming Huang, Wei Chu, Fangzhen Lin, Wenhu Chen
(**Spotlight, Top 10% of accepted papers**)

[**NeurIPS 2025**] MoCha: Towards Movie-Grade Talking Character Synthesis
Cong Wei, Bo Sun, Haoyu Ma, Ji Hou, Felix Juefei-Xu, Zecheng He, Xiaoliang Dai, Luxin Zhang, Kunpeng Li, Tingbo Hou, Animesh Sinha, Peter Vajda, Wenhu Chen
(**Spotlight, Top 10% of accepted papers**)

[**TMLR 2025**] PixelWorld: Towards Perceiving Everything as Pixels
Zhiheng Lyu, Xueguang Ma, Wenhu Chen

[**Findings of EMNLP 2025**] VisCoder: Fine-Tuning LLMs for Executable Python Visualization Code Generation
Yuansheng Ni, Ping Nie, Kai Zou, Xiang Yue, Wenhu Chen

[**EMNLP 2025**] VisualWebInstruct: Scaling up Multimodal Instruction Data through Web Search
Yiming Jia, Jiachen Li, Xiang Yue, Bo Li, Ping Nie, Kai Zou, Wenhu Chen

[**EMNLP 2025**] Unleashing the Reasoning Potential of Pre-trained LLMs by Critique Fine-Tuning on One Problem
Yubo Wang, Ping Nie, Kai Zou, Lijun Wu, Wenhu Chen

[**LLM Report 2025**] Gemini 2.5: Pushing the frontier with advanced reasoning, multimodality, long context, and next generation agentic capabilities
Gemini Team

[**TMLR 2025**] ABC: Achieving Better Control of Multimodal Embeddings using VLMs
Benjamin Schneider, Florian Kerschbaum, Wenhu Chen

[**COLM 2025**] ScholarCopilot: Training Large Language Models for Academic Writing with Accurate Citations
Yubo Wang, Xueguang Ma, Ping Nie, Huaye Zeng, Zhiheng Lyu, Yuxuan Zhang, Benjamin Schneider, Yi Lu, Xiang Yue, Wenhu Chen

[**COLM 2025**] Critique Fine-Tuning: Learning to Critique is More Effective than Learning to Imitate
Yubo Wang, Xiang Yue, Wenhu Chen

[**ICCV 2025**] Vamba: Understanding Hour-Long Videos with Hybrid Mamba-Transformers
Weiming Ren, Wentao Ma, Huan Yang, Cong Wei, Ge Zhang, Wenhu Chen

[**ACL 2025**] AceCoder: Acing Coder RL via Automated Test-Case Synthesis
Huaye Zeng, Dongfu Jiang, Haozhe Wang, Ping Nie, Xiaotong Chen, Wenhu Chen

[**ACL 2025**] TheoremExplainAgent: Towards Video-based Multimodal Explanations for LLM Theorem Understanding
Max Ku, Thomas Chong, Jonathan Leung, Krish Shah, Alvin Yu, Wenhu Chen

[**CVPR 2025**] Enhancing Long-Duration and High-Resolution Video Understanding by Video SpatioTemporal Augmentation
Weiming Ren, Huan Yang, Jie Min, Cong Wei, Wenhu Chen

[**TMLR 2025**] Long-context LLMs Struggle with Long In-context Learning
Tianle Li, Ge Zhang, Quy Duc Do, Xiang Yue, Wenhu Chen

[**ICLR 2025**] OmniEdit: Building Image Editing Generalist Models Through Specialist Supervision
Cong Wei, Zheyang Xiong, Weiming Ren, Xinrun Du, Ge Zhang, Wenhu Chen

[**ICLR 2025**] VLM2Vec: Training Vision-Language Models for Massive Multimodal Embedding Tasks
Ziyan Jiang, Rui Meng, Xinyi Yang, Semih Yavuz, Yingbo Zhou, Wenhu Chen

[**ICLR 2025**] MEGA-Bench: Scaling Multimodal Evaluation to over 500 Real-World Tasks
Jiacheng Chen, Tianhao Liang, Sherman Siu, Zhengqing Wang, Kai Wang, Yubo Wang, Yuansheng Ni, Wang Zhu, Ziyan Jiang, Bohan Lyu, Dongfu Jiang, Xuan He, Yuan Liu, Hexiang Hu, Xiang Yue, Wenhu Chen

[**TMLR 2024**] MANTIS: Interleaved Multi-Image Instruction Tuning
Dongfu Jiang, Xuan He, Huaye Zeng, Cong Wei, Max Ku, Qian Liu, Wenhu Chen
(Best Paper Award)

[**TMLR 2024**] AnyV2V: A Plug-and-Play Framework For Any Video-to-Video Editing Tasks
Max Ku, Cong Wei, Weiming Ren, Huan Yang, Wenhu Chen
(Reproducibility Certification)

[**NeurIPS 2024**] MMLU-Pro: A More Robust and Challenging Multi-Task Language Understanding Benchmark
Yubo Wang, Xueguang Ma, Ge Zhang, Yuansheng Ni, Abhranil Chandra, Shiguang Guo, Weiming Ren, Aaran Arulraj, Xuan He, Ziyan Jiang, Tianle Li, Max Ku, Kai Wang, Alex Zhuang, Rongqi Fan, Xiang Yue, Wenhu Chen
(Spotlight, Top 10% of accepted papers)

[**NeurIPS 2024**] GenAI Arena: An Open Evaluation Platform for Generative Models
Dongfu Jiang, Max Ku, Tianle Li, Yuansheng Ni, Shizhuo Sun, Rongqi Fan, Wenhu Chen

[**NeurIPS 2024**] MAMmoTH2: Scaling Instructions from the Web
Xiang Yue, Tuney Zheng, Ge Zhang, Wenhu Chen

[**EMNLP 2024**] Unifying Multimodal Retrieval via Document Screenshot Embedding
Xueguang Ma, Sheng-Chieh Lin, Minghan Li, Wenhu Chen, Jimmy Lin

[**EMNLP 2024**] VideoScore: Building Automatic Metrics to Simulate Fine-grained Human Feedback for Video Generation

Xuan He, Dongfu Jiang, Ge Zhang, Max Ku, Achint Soni, Sherman Siu, Haonan Chen, Abhuranil Chandra, Ziyang Jiang, Aaran Arulraj, Kai Wang, Quy Duc Do, Yuansheng Ni, Bohan Lyu, Yaswanth Narsupalli, Rongqi Fan, Zhiheng Lyu, Yuchen Lin, Wenhui Chen

[**Findings of EMNLP 2024**] Augmenting Black-box LLMs with Medical Textbooks for Clinical Question Answering
Yubo Wang, Xueguang Ma, Wenhui Chen

[**COLM 2024**] StructLM: Towards Building Generalist Models for Structured Knowledge Grounding

Alex Zhuang, Ge Zhang, Tianyu Zheng, Xinrun Du, Junjie Wang, Weiming Ren, Stephen W. Huang, Jie Fu, Xiang Yue, Wenhui Chen

[**ACL 2024**] VIEScore: Towards Explainable Metrics for Conditional Image Synthesis Evaluation

Max Ku, Dongfu Jiang, Cong Wei, Xiang Yue, Wenhui Chen

[**Findings of ACL 2024**] OpenCodeInterpreter: Integrating Code Generation with Execution and Refinement

Tianyu Zheng, Ge Zhang, Tianhao Shen, Xueling Liu, Bill Yuchen Lin, Jie Fu, Wenhui Chen, Xiang Yue

[**TMLR 2024**] TIGERScore: Towards Building Explainable Metric for All Text Generation Tasks

Dongfu Jiang*, Yishan Li*, Ge Zhang, Wenhao Huang, Bill Yuchen Lin, Wenhui Chen

[**ECCV 2024**] UniIR: Training and Benchmarking Universal Multimodal Information Retrievers

Cong Wei, Yang Chen, Haonan Chen, Hexiang Hu, Ge Zhang, Jie Fu, Alan Ritter, Wenhui Chen

(Oral, Top 5% of accepted papers)

[**CVPR 2024**] Instruct-Imagen: Image Generation with Multi-modal Instruction

Hexiang Hu*, Kelvin C.K. Chan*, Yu-Chuan Su*, Wenhui Chen*, Yandong Li, Kihyuk Sohn, Yang Zhao, Xue Ben, Boqing Gong, William Cohen, Ming-Wei Chang, Xuhui Jia

(Oral, Top 3% of accepted papers)

[**CVPR 2024**] MMMU: A Massive Multi-discipline Multimodal Understanding and Reasoning Benchmark for Expert AGI

Xiang Yue, Yuansheng Ni, Kai Zhang, Tianyu Zheng, Ruqi Liu, Ge Zhang, Samuel Stevens, Dongfu Jiang, Weiming Ren, Yuxuan Sun, Cong Wei, Botao Yu, Ruibin Yuan, Renliang Sun, Ming Yin, Boyuan Zheng, Zhenzhu Yang, Yibo Liu, Wenhao Huang, Huan Sun, Yu Su, Wenhui Chen

(Best Paper Finalist, Top 1% of accepted papers)

[**ICLR 2024**] ImagenHub: Standardizing the evaluation of conditional image generation models

Max Ku, Tianle Li, Kai Zhang, Yujie Lu, Xingyu Fu, Wenwen Zhuang, Wenhui Chen

[**ICLR 2024**] MAMMO-TH: Building Math Generalist Models through Hybrid Instruction Tuning

Xiang Yue*, Xingwei Qu, Ge Zhang, Yao Fu, Wenhao Huang, Huan Sun, Yu Su, Wenhui Chen*

(Spotlight, Top 20% of accepted papers)

[**WACV 2024**] Synthesizing Coherent Story with Auto-Regressive Latent Diffusion Models

Xichen Pan, Pengda Qin, Yuhong Li, Hui Xue, Wenhui Chen

(Oral, Top 5% of accepted papers)

[**TMLR 2023**] DreamEdit: Subject-driven Image Editing

Tianle Li, Max Ku, Cong Wei, Wenhui Chen

[**TMLR 2023**] Program of Thoughts Prompting: Disentangling Computation from Reasoning for Numerical Reasoning Tasks

Wenhui Chen*, Xueguang Ma*, Xinyi Wang, William W. Cohen

[**EMNLP 2023**] TheoremQA: A Theorem-driven Question Answering dataset

Wenhui Chen, Ming Yin, Max Ku, Elaine Wan, Xueguang Ma, Jianyu Xu, Tony Xia, Xinyi Wang, Pan Lu

[**NeurIPS 2023**] Subject-driven Text-to-Image Generation via Apprenticeship Learning

Wenhui Chen, Hexiang Hu, Yandong Li, Nataniel Ruiz, Xuhui Jia, Ming-Wei Chang, William W. Cohen

[**ACL 2023**] Few-shot In-context Learning on Knowledge Base Question Answering

Tianle Li, Xueguang Ma, Alex Zhuang, Yu Gu, Yu Su and Wenhui Chen

[**Findings of EACL 2023**] Large Language Models are few(1)-shot Table Reasoners

Wenhu Chen

[**EACL 2023**] Augmenting Pre-trained Language Models with QA-Memory for Open-Domain Question Answering

Wenhu Chen, Pat Verga, Michiel de Jong, John Wieting, William Cohen

[**ICLR 2023**] Re-Imagen: Retrieval-Augmented Text-to-Image Generator

Wenhu Chen, Hexiang Hu, Chitwan Saharia, William W. Cohen

[**EMNLP 2022**] MuRAG: Multimodal Retrieval-Augmented Generator for Open Question Answering over Images and Text

Wenhu Chen, Hexiang Hu, Xi Chen, Pat Verga, William W. Cohen

[**NeurIPS 2021**] A Dataset for Answering Time-Sensitive Questions

Wenhu Chen, Xinyi Wang, William Yang Wang

[**EMNLP 2020**] KGPT: Knowledge-Grounded Pre-Training for Data-to-Text Generation

Wenhu Chen, Yu Su, Xifeng Yan, William Yang Wang.

[**Findings of EMNLP 2020**] HybridQA: A Dataset of Multi-Hop Question Answering over Tabular and Textual Data

Wenhu Chen, Hanwen Zha, Zhiyu Chen, Wenhan Xiong, Hong Wang, William Wang.

[**WACV 2021**] Meta Module Network for Compositional Visual Reasoning

Wenhu Chen, Zhe Gan, Linjie Li, Yu Cheng, William Wang, Jingjing Liu.

(Best Student Paper Honorable Mention, Top 1% of accepted papers)

[**Findings of EMNLP 2020**] HybridQA: A Dataset of Multi-Hop Question Answering over Tabular and Textual Data

Wenhu Chen, Hanwen Zha, Zhiyu Chen, Wenhan Xiong, Hong Wang, William Wang.

[**ACL 2020**] Logical Natural Language Generation from Open-Domain Tables

Wenhu Chen, Jianshu Chen, Yu Su, Zhiyu Chen and William Yang Wang.

[**ICLR 2020**] TabFact: A Large-scale Dataset for Table-based Fact Verification

Wenhu Chen, Hongmin Wang, Jianshu Chen, Yunkai Zhang, Hong Wang, Shiyang Li, Xiyu Zhou and William Yang Wang.

[**ACL 2019**] Semantically Conditioned Dialog Response Generation via Hierarchical Disentangled Self-Attention

Wenhu Chen, Jianshu Chen, Pengda Qin, Xifeng Yan and William Yang Wang.

[**NAACL 2019**] How Large A Vocabulary Does Text Classification Need? A Variational Approach on Vocabulary Selection

Wenhu Chen, Yu Su, Yilin Shen, Zhiyu Chen, Xifeng Yan and William Yang Wang.

[**EMNLP 2018**] XL-NBT: A Cross-lingual Neural Belief Tracking Framework

Wenhu Chen, Jianshu Chen, Yu Su, Xin Wang, Dong Yu, Xifeng Yan and William Yang Wang.

[**ACL 2018**] No Metrics Are Perfect: Adversarial Reward Learning for Visual Storytelling

Wenhu Chen*, Xin Wang*, Yuan-Fang Wang and William Yang Wang.

[**NAACL 2018**] Variational Knowledge Graph Reasoning

Wenhu Chen, Wenhan Xiong, William Yang Wang, Xifeng Yan.

[**NAACL 2018**] Generative Bridging Network in Neural Sequence Prediction

Wenhu Chen, Guanlin Li, Shuo Ren, Shujie Liu, Zhirui Zhang, Mu Li, Ming Zhou.

ACADEMIC AWARD

2025 TMLR Outstanding Paper Award.

2025 Math Golden Jubilee Award.

2024 CVPR Best Paper Finalist.

2023 AACL-IJCNLP Area Chair Award.

2022 CIFAR AI Chair.

2021 WACV Best Student Paper Honorable Mention.

2019 NeurIPS Outstanding Reviewer Award.

FUNDINGS

CIFAR Solution Network: 2026 - 2028

JELF-CFI Award: 2025 - 2030

CIFAR AI Chair Fund: 2022 - 2027

NSERC Discovery Fund: 2023 - 2028

Mitacs Accelerate Fund: 2023 - 2025

CIFAR AI Catalyst Fund: 2024 - 2026

Canada National Research Council (AI4D): 2024 - 2026

Canada National Research Council (New Beginning): 2025 - 2026

SERVICE

Program Committee: NAACL, ACL, EMNLP, ICML, NeurIPS, ICLR

Area Chair: AAAI, EMNLP, ACL, NAACL